

How to Jump-Start Your Deep Learning Research

Florian Knoll

Department of Radiology, CBI
Center for Advanced Imaging Innovation and Research (CAI²R)
NYU School of Medicine, New York, USA



JOINT ANNUAL MEETING
ISMRM-ESMRMB
16–21 June 2018

SMRT 27th Annual Meeting 15–18 June 2018
www.smrt.org

Paris Expo Porte de Versailles
Paris, France

Declaration of Financial Interests or Relationships

Speaker Name: Florian Knoll

I have no financial interests or relationships to disclose with regard to the subject matter of this presentation.

Surfing the big AI wave



Getting started...



Outline

- Didactic implementation examples
 - Linear regression
 - Classification of DTI data (MLP)
 - Classification of reconstructed image quality (CNN)

```
for ii in range(training_epochs):
    for x_train_batch, y_train_batch in zip(x_train,y_train):
        loss = loss_fcn(x_train_batch, y_train_batch)
        loss_ii[ii] = loss.item()
        k_ii[ii] = k.data
        d_ii[ii] = d.data
        loss.backward()
        k.data = k.data - lr * k.grad.data
        d.data = d.data - lr * d.grad.data
```

Outline

- Didactic implementation examples
 - Linear regression
 - Classification of DTI data (MLP)
 - Classification of reconstructed image quality (CNN)

```
for ii in range(training_epochs):
    for x_train_batch, y_train_batch in zip(x_train,y_train):
        loss = loss_fcn(x_train_batch, y_train_batch)
        loss_ii[ii] = loss.item()
        k_ii[ii] = k.data
        d_ii[ii] = d.data
        loss.backward()
        k.data = k.data - lr * k.grad.data
        d.data = d.data - lr * d.grad.data
```

The screenshot shows a GitHub repository page for 'ISMRM2018_Educational_DeepLearning'. The top navigation bar includes '2 commits', '1 branch', '0 releases', '1 contributor', and 'MIT' license information. Below the navigation bar, there are buttons for 'Branch: master', 'New pull request', 'Create new file', 'Upload files', 'Find file', and 'Clone or download'. The main content area displays a list of files in the repository:

File	Description	Last Commit
CNN_recon_quality_classification_tensorflow.py	Initial check in for ISMRM	16 hours ago
LICENSE	Initial commit	16 hours ago
README.md	Initial check in for ISMRM	16 hours ago
data.zip	Initial check in for ISMRM	16 hours ago
dti_classification_keras.py	Initial check in for ISMRM	16 hours ago
dti_classification_matlab.m	Initial check in for ISMRM	16 hours ago
dti_classification_pytorch.py	Initial check in for ISMRM	16 hours ago
linear_regression_pytorch.py	Initial check in for ISMRM	16 hours ago
linear_regression_pytorch_cuda.py	Initial check in for ISMRM	16 hours ago
linear_regression_tensorflow.py	Initial check in for ISMRM	16 hours ago
rgb2gray.py	Initial check in for ISMRM	16 hours ago

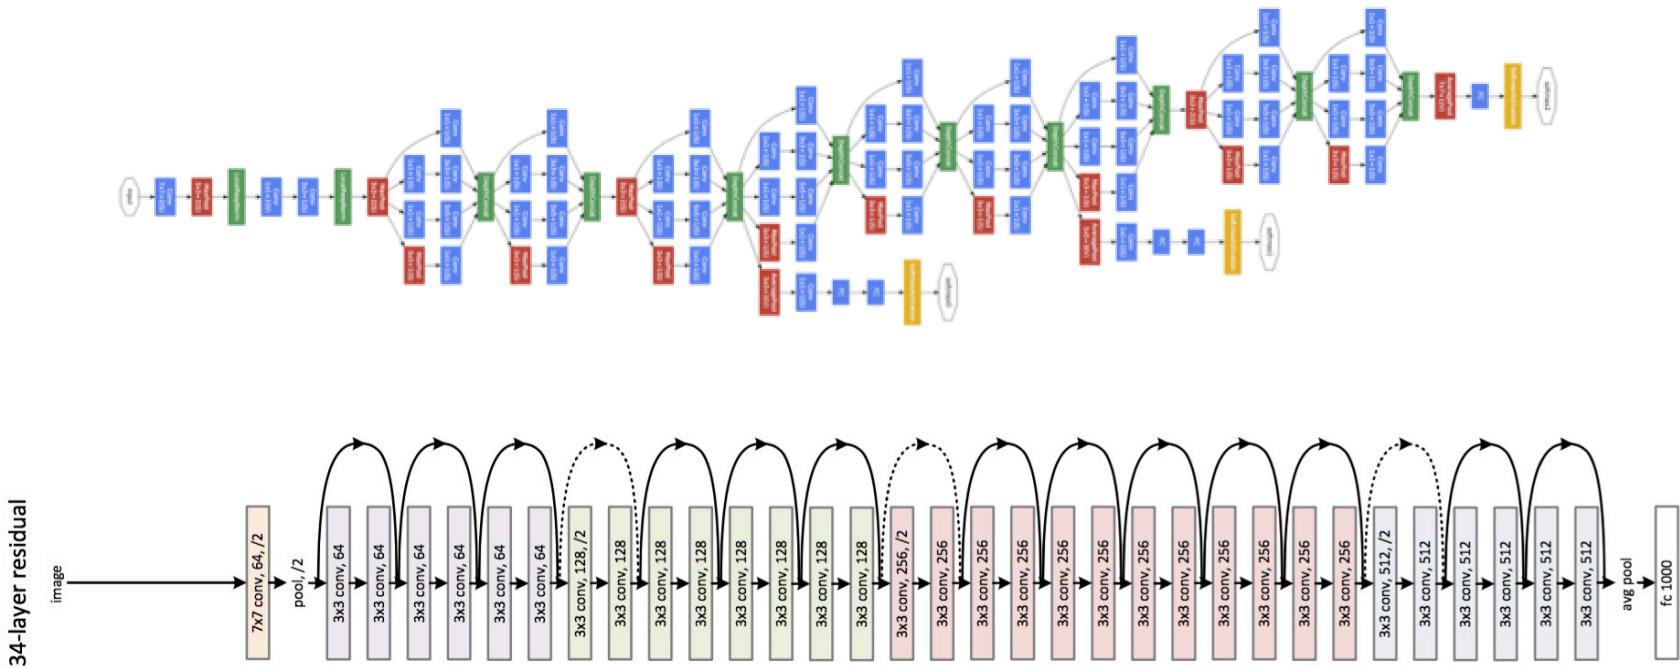
Outline

- Didactic implementation examples
 - Linear regression
 - Classification of DTI data (MLP)
 - Classification of reconstructed image quality (CNN)
- Software frameworks
- Generalization: Model complexity, overfitting vs. underfitting
- Architectures: Convolutional vs fully connected

```
for ii in range(training_epochs):
    for x_train_batch, y_train_batch in zip(x_train,y_train):
        loss = loss_fcn(x_train_batch, y_train_batch)
        loss_i[ii] = loss.item()
        k_i[ii] = k.data
        d_i[ii] = d.data
        loss.backward()
        k.data = k.data - lr * k.grad.data
        d.data = d.data - lr * d.grad.data
```

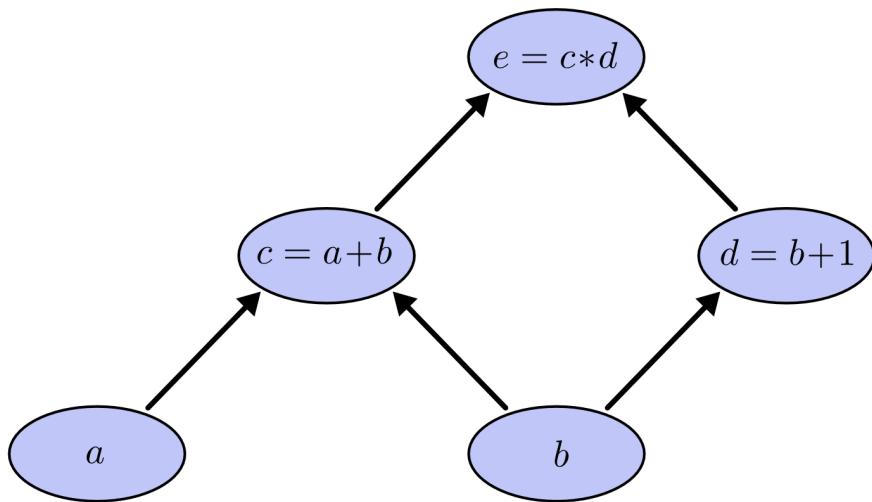
Software frameworks:
Why would you want to use one?

Deep learning models



Szegedy et al., CVPR 2015
He et al., CVPR 2016

Computational graphs



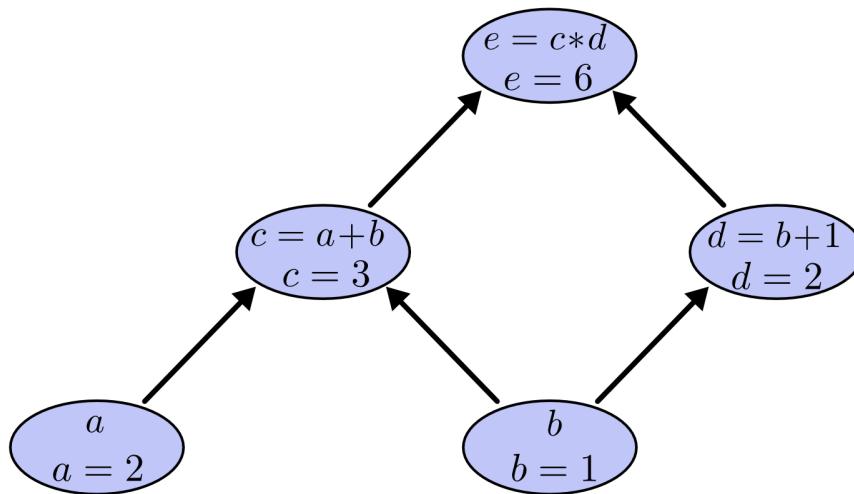
$$e = (a + b)(b + 1)$$

$$c = a + b$$

$$d = b + 1$$

$$e = c \cdot d$$

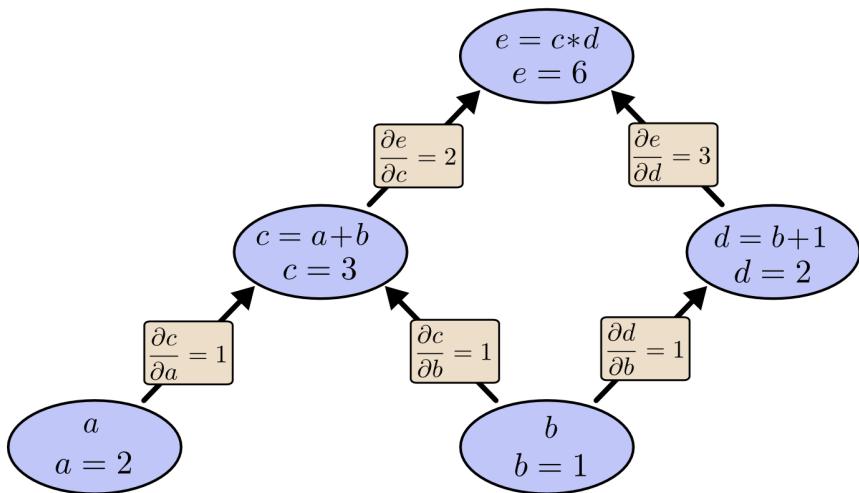
Computational graphs



$$a = 2$$

$$b = 1$$

Computational graphs



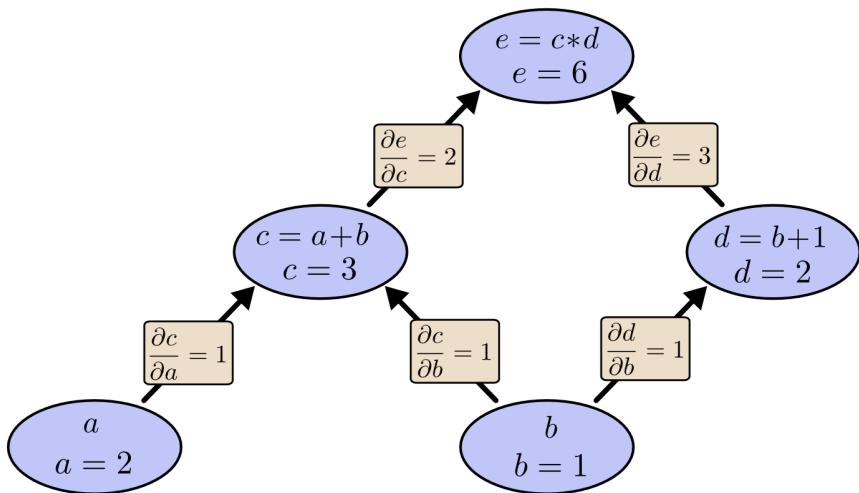
$$c = a + b$$

$$\frac{\partial c}{\partial a} = b$$

$$e = c \cdot d$$

$$\frac{\partial e}{\partial c} = d$$

Computational graphs



$$c = a + b$$

$$\frac{\partial c}{\partial a} = b$$

$$e = c \cdot d$$

$$\frac{\partial e}{\partial c} = d$$

Data and computation



<http://www.image-net.org/>

SKYNET.NYUMC.ORG

[knollf01@skynet ~]\$ bhosts	HOST_NAME	STATUS	JL/U	MAX	NJOBS	RUN	SSUSP	USUSP	RSV
skygpu01	closed	-	80	0	0	0	0	0	0
skygpu02	ok	-	80	0	0	0	0	0	0
skygpu03	closed	-	80	0	0	0	0	0	0
skygpu04	closed	-	80	0	0	0	0	0	0
skygpu05	Closed	-	80	0	0	0	0	0	0
skygpu06	ok	-	80	20	20	0	0	0	0
skygpu07	ok	-	80	0	0	0	0	0	0
skygpu08	ok	-	80	0	0	0	0	0	0
skygpu09	ok	-	80	0	0	0	0	0	0
skygpu10	ok	-	80	0	0	0	0	0	0
skygpu11	Closed	-	160	0	0	0	0	0	0
skygpu12	Closed	-	160	0	0	0	0	0	0
skygpu13	ok	-	80	0	0	0	0	0	0
skygpu14	Closed	-	80	0	0	0	0	0	0
skygpu15	Closed	-	80	0	0	0	0	0	0
skygpu16	ok	-	40	16	16	0	0	0	0
skygpu17	ok	-	80	0	0	0	0	0	0
skynet	Closed	-	40	0	0	0	0	0	0
skynet2	Closed	-	80	0	0	0	0	0	0



Functionalities of software frameworks

1. Data structures for (large) computational graphs
2. Automatic calculation of gradients
3. GPU support
4. High level support for architecture design
5. Data handling functions (minibatches etc.)
6. Numerical optimizers

Popularity of deep learning frameworks

Library	Rank	Overall	Github	Stack Overflow	Google Results
tensorflow	1	10.87	4.25	4.37	2.24
keras	2	1.93	0.61	0.83	0.48
caffe	3	1.86	1.00	0.30	0.55
theano	4	0.76	-0.16	0.36	0.55
pytorch	5	0.48	-0.20	-0.30	0.98
sonnet	6	0.43	-0.33	-0.36	1.12
mxnet	7	0.10	0.12	-0.31	0.28
torch	8	0.01	-0.15	-0.01	0.17
cntk	9	-0.02	0.10	-0.28	0.17
dlib	10	-0.60	-0.40	-0.22	0.02

September 2017: <https://github.com/thedataincubator/data-science-blogs>

ISMRM, June 16th 2018

First generation deep learning frameworks

theano



Caffe



Second generation deep learning frameworks



TensorFlow

Google

P Y T H O N



K Keras

m xnet

amazon



Microsoft

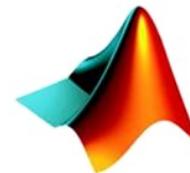
Second generation deep learning frameworks



TensorFlow

PYTORCH

K Keras

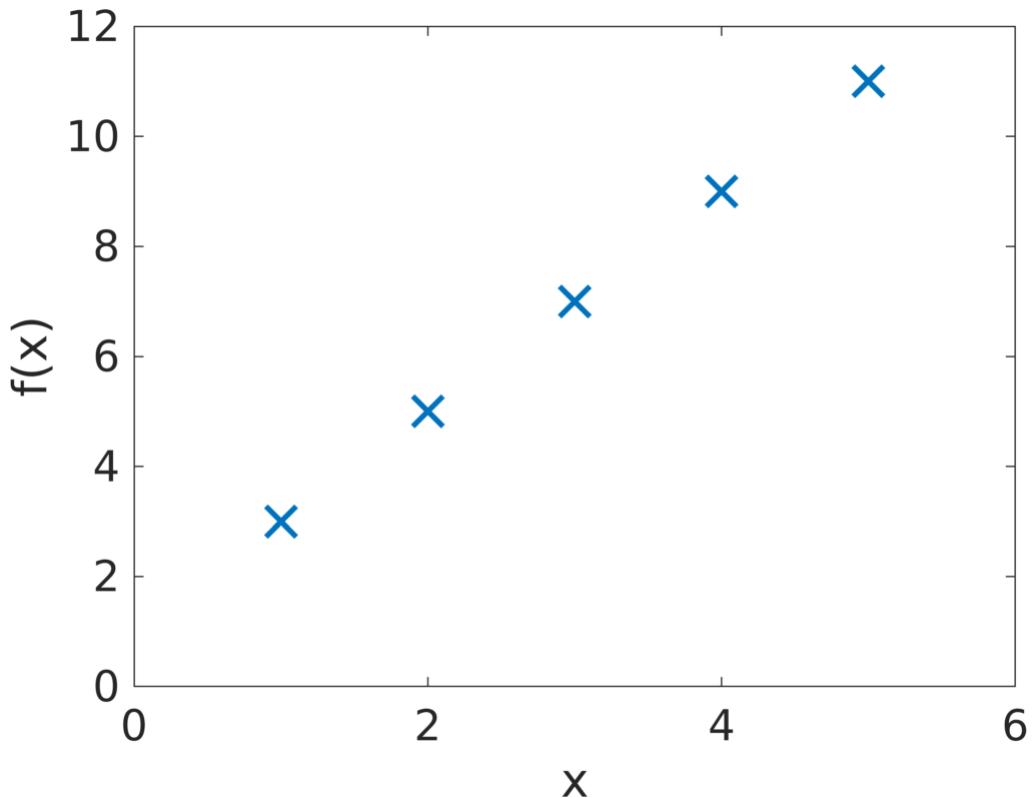


MATLAB®

Example 1: Linear Regression

A low level view of PyTorch and TensorFlow

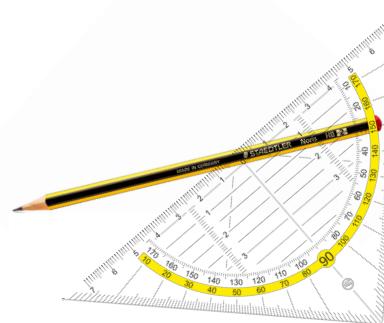
Data



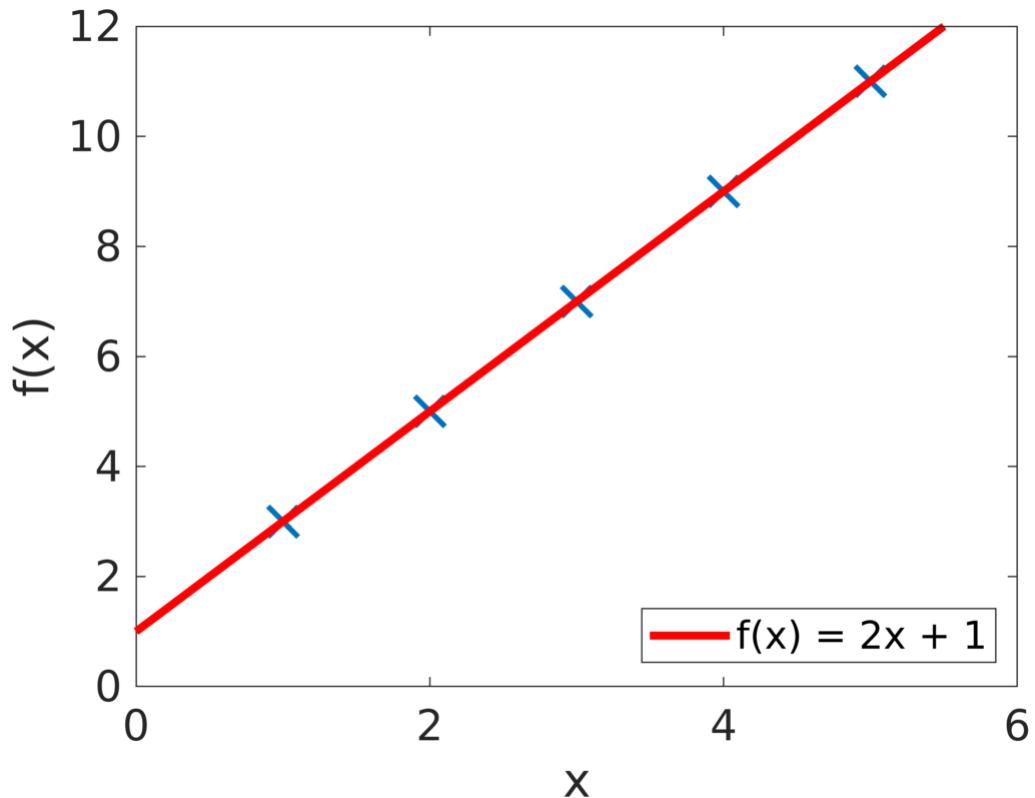
$$x = [1, 2, 3, 4, 5]$$

$$y = [3, 5, 7, 9, 11]$$

$$f(x) = kx + d$$



Data



$$x = [1, 2, 3, 4, 5]$$

$$y = [3, 5, 7, 9, 11]$$

$$f(x) = kx + d$$

$$f(x) = 2x + 1$$

Linear regression with gradient descent: PyTorch

```
## training data
x_train = torch.Tensor([1,2,3,4,5]).float()
y_train = torch.Tensor([3,5,7,9,11]).float()

## model parameters
k = torch.Tensor([0.1]).float()
k.requires_grad=True

d = torch.Tensor([-0.1]).float()
d.requires_grad=True

## model and optimizer
def forward(x):
    return x * k + d

def loss_fcn(x,y):
    y_pred = forward(x)
    return (y_pred-y) * (y_pred-y)

## train
training_epochs = 1000
lr = 0.005
loss_ii = np.zeros(training_epochs)
k_ii = np.zeros(training_epochs)
d_ii = np.zeros(training_epochs)

t = time.time()
for ii in range(training_epochs):
    for x_train_batch, y_train_batch in zip(x_train,y_train):
        loss = loss_fcn(x_train_batch, y_train_batch)
        loss_ii[ii] = loss.item()
        k_ii[ii] = k.data
        d_ii[ii] = d.data
        loss.backward()
        k.data = k.data - lr * k.grad.data
        d.data = d.data - lr * d.grad.data

    # Manually set gradient to zero after update step to prevent gradient accumulation
    k.grad.data.zero_()
    d.grad.data.zero_()
    print('epoch: {}, k={:.3}, d={:.3}, loss={:.3}'.format(ii+1,k_ii[ii], d_ii[ii], loss_ii[ii]))

elapsed = time.time() - t
print('Training time: {:.2} s'.format(elapsed))
```



https://github.com/FlorianKnoll/ISMRM2018_Educational_DeepLearning

Training data and model parameters

```
## training data
x_train = torch.Tensor([1,2,3,4,5]).float()
y_train = torch.Tensor([3,5,7,9,11]).float()
```

```
## model parameters
k = torch.Tensor([0.1]).float()
k.requires_grad=True
```

```
d = torch.Tensor([-0.1]).float()
d.requires_grad=True
```



Automatic differentiation

Model and loss function

```
## training data
x_train = torch.Tensor([1,2,3,4,5]).float()
y_train = torch.Tensor([3,5,7,9,11]).float()

## model parameters
k = torch.Tensor([0.1]).float()
k.requires_grad=True

d = torch.Tensor([-0.1]).float()
d.requires_grad=True

## model and optimizer
def forward(x):
    return x * k + d

def loss_fcn(x,y):
    y_pred = forward(x)
    return (y_pred-y) * (y_pred-y)

## train
training_epochs = 1000
lr = 0.005
loss_ii = np.zeros(training_epochs)
k_ii = np.zeros(training_epochs)
d_ii = np.zeros(training_epochs)

t = time.time()
for ii in range(training_epochs):
    for x_train_batch, y_train_batch in zip(x_train,y_train):
        loss = loss_fcn(x_train_batch, y_train_batch)
        loss_ii[ii] = loss.item()
        k_ii[ii] = k.data
        d_ii[ii] = d.data
        loss.backward()
        k.data = k.data - lr * k.grad.data
        d.data = d.data - lr * d.grad.data

    # Manually set gradient to zero after update step to prevent gradient accumulation
    k.grad.data.zero_()
    d.grad.data.zero_()
    print('epoch: {}, k={:.3}, d={:.3}, loss={:.3}'.format(ii+1,k_ii[ii], d_ii[ii], loss_ii[ii]))

elapsed = time.time() - t
print('Training time: {:.2} s'.format(elapsed))
```



Training the model

```
## training data
x_train = torch.Tensor([1,2,3,4,5]).float()
y_train = torch.Tensor([3,5,7,9,11]).float()

## model parameters
k = torch.Tensor([0.1]).float()
k.requires_grad=True

d = torch.Tensor([-0.1]).float()
d.requires_grad=True

## model and optimizer
def forward(x):
    return x * k + d

def loss_fcn(x,y):
    y_pred = forward(x)
    return (y_pred-y) * (y_pred-y)

## train
training_epochs = 1000
lr = 0.005
loss_ii = np.zeros(training_epochs)
k_ii = np.zeros(training_epochs)
d_ii = np.zeros(training_epochs)

t = time.time()
for ii in range(training_epochs):
    for x_train_batch, y_train_batch in zip(x_train,y_train):
        loss = loss_fcn(x_train_batch, y_train_batch)
        loss_ii[ii] = loss.item()
        k_ii[ii] = k.data
        d_ii[ii] = d.data
        loss.backward()
        k.data = k.data - lr * k.grad.data
        d.data = d.data - lr * d.grad.data

        # Manually set gradient to zero after update step to prevent gradient accumulation
        k.grad.data.zero_()
        d.grad.data.zero_()
    print('epoch: {}, k={:.3}, d={:.3}, loss={:.3}'.format(ii+1,k_ii[ii], d_ii[ii], loss_ii[ii]))

elapsed = time.time() - t
print('Training time: {:.2} s'.format(elapsed))
```



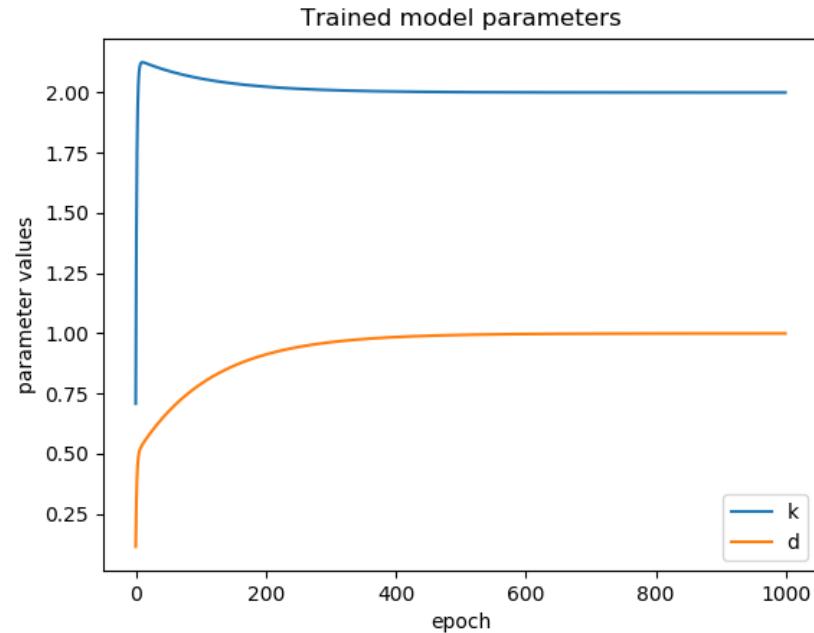
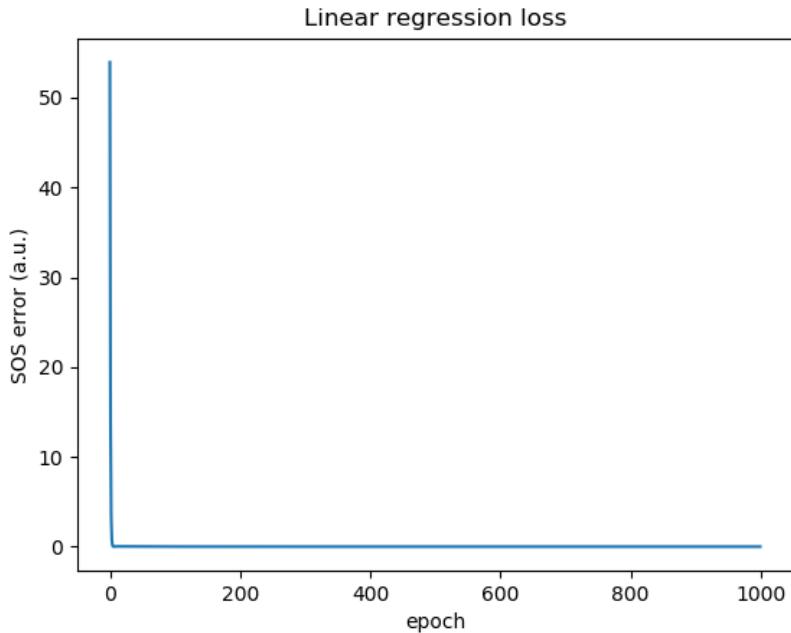
Loop over samples, update parameters

```
for ii in range(training_epochs):
    for x_train_batch, y_train_batch in zip(x_train,y_train):
        loss = loss_fcn(x_train_batch, y_train_batch)
        loss_ii[ii] = loss.item()
        k_ii[ii] = k.data
        d_ii[ii] = d.data
        loss.backward()
        k.data = k.data - lr * k.grad.data
        d.data = d.data - lr * d.grad.data
```



The computational graph is dynamically created during the iterations

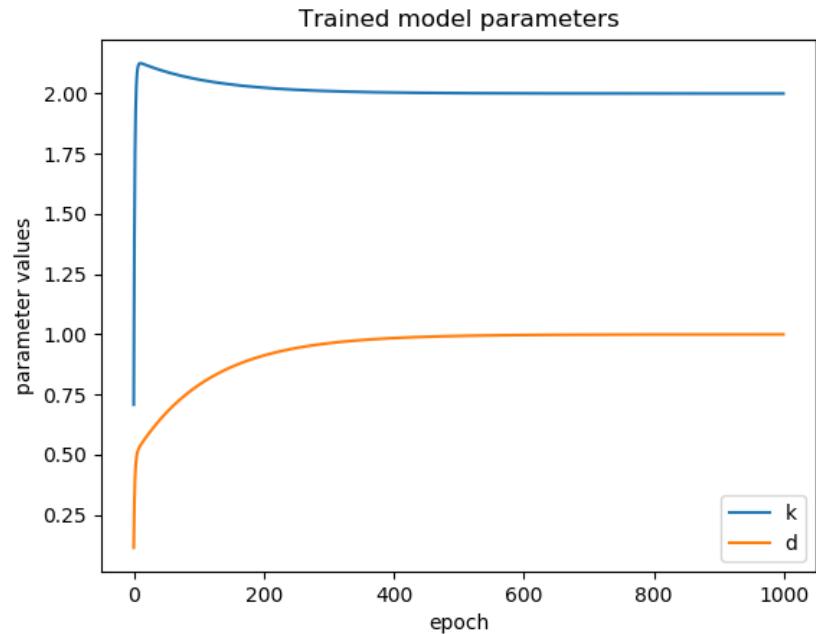
Linear regression with gradient descent: PyTorch



Linear regression in PyTorch

```
epoch: 993, k=2.0, d=1.0, loss=1.11e-09
epoch: 994, k=2.0, d=1.0, loss=1.11e-09
epoch: 995, k=2.0, d=1.0, loss=9.9e-10
epoch: 996, k=2.0, d=1.0, loss=9.9e-10
epoch: 997, k=2.0, d=1.0, loss=9.9e-10
epoch: 998, k=2.0, d=1.0, loss=9.9e-10
epoch: 999, k=2.0, d=1.0, loss=9.31e-10
epoch: 1000, k=2.0, d=1.0, loss=9.9e-10
Training time: 1.4 s
```

$$f(x) = 2x + 1$$



Linear regression in PyTorch on the CPU

```
## training data
x_train = torch.Tensor([1,2,3,4,5]).float()
y_train = torch.Tensor([3,5,7,9,11]).float()



---


## model parameters
k = torch.Tensor([0.1]).float()
k.requires_grad=True

d = torch.Tensor([-0.1]).float()
d.requires_grad=True
```

Linear regression in PyTorch on the GPU

```
## training data
x_train = torch.Tensor([1,2,3,4,5]).float().to(torch.device("cuda"))
y_train = torch.Tensor([3,5,7,9,11]).float().to(torch.device("cuda"))



---


## model parameters
k = torch.Tensor([0.1]).float().to(torch.device("cuda"))
k.requires_grad=True

d = torch.Tensor([-0.1]).float().to(torch.device("cuda"))
d.requires_grad=True
```

Linear regression in TensorFlow

```
## training data
x_train = [1,2,3,4,5]
y_train = [3,5,7,9,11]

## Model parameters
k = tf.Variable([0.1], dtype=tf.float32)
d = tf.Variable([-0.1], dtype=tf.float32)

# Model input and output
x = tf.placeholder(tf.float32)
linear_model = (k * x + d)
y = tf.placeholder(tf.float32)

# sum of squares loss
loss = tf.reduce_sum(tf.square(linear_model - y))

## optimizer
lr = 0.005
optimizer = tf.train.GradientDescentOptimizer(lr).minimize(loss)

## train
training_epochs = 1000
init = tf.global_variables_initializer()
sess = tf.Session()
sess.run(init)
loss_ii = np.zeros(training_epochs)
k_ii = np.zeros(training_epochs)
d_ii = np.zeros(training_epochs)

t = time.time()
for ii in range(training_epochs):
    [k_ii_temp, d_ii_temp, loss_ii[ii]] = sess.run([k, d, loss], {x:x_train, y:y_train})
    k_ii[ii] = k_ii_temp.item()
    d_ii[ii] = d_ii_temp.item()
    print('epoch: {}, k={:.3}, b={:.3}, loss={:.3}'.format(ii+1,k_ii[ii], d_ii[ii], loss_ii[ii]))
    sess.run(optimizer, {x:x_train, y:y_train})

elapsed = time.time() - t
print('Training time: {:.2} s'.format(elapsed))
```



Static computational graph

Linear regression in TensorFlow

```
## training data
x_train = [1,2,3,4,5]
y_train = [3,5,7,9,11]

## Model parameters
k = tf.Variable([0.1], dtype=tf.float32)
d = tf.Variable([-0.1], dtype=tf.float32)

# Model input and output
x = tf.placeholder(tf.float32)
linear_model = (k * x + d)
y = tf.placeholder(tf.float32)

# sum of squares loss
loss = tf.reduce_sum(tf.square(linear_model - y))
```



Linear regression in TensorFlow

```
## training data
x_train = [1,2,3,4,5]
y_train = [3,5,7,9,11]

## Model parameters
k = tf.Variable([0.1], dtype=tf.float32)
d = tf.Variable([-0.1], dtype=tf.float32)

# Model input and output
x = tf.placeholder(tf.float32)
linear_model = (k * x + d)
y = tf.placeholder(tf.float32)

# sum of squares loss
loss = tf.reduce_sum(tf.square(linear_model - y))

## optimizer
lr = 0.005
optimizer = tf.train.GradientDescentOptimizer(lr).minimize(loss)
```



Linear regression in TensorFlow

```
## training data
x_train = [1,2,3,4,5]
y_train = [3,5,7,9,11]

## Model parameters
k = tf.Variable([0.1], dtype=tf.float32)
d = tf.Variable([-0.1], dtype=tf.float32)

# Model input and output
x = tf.placeholder(tf.float32)
linear_model = (k * x + d)
y = tf.placeholder(tf.float32)

# sum of squares loss
loss = tf.reduce_sum(tf.square(linear_model - y))

## optimizer
lr = 0.005
optimizer = tf.train.GradientDescentOptimizer(lr).minimize(loss)

## train
training_epochs = 1000
init = tf.global_variables_initializer()
sess = tf.Session()
sess.run(init)
loss_hi = np.zeros(training_epochs)
k_hi = np.zeros(training_epochs)
d_hi = np.zeros(training_epochs)

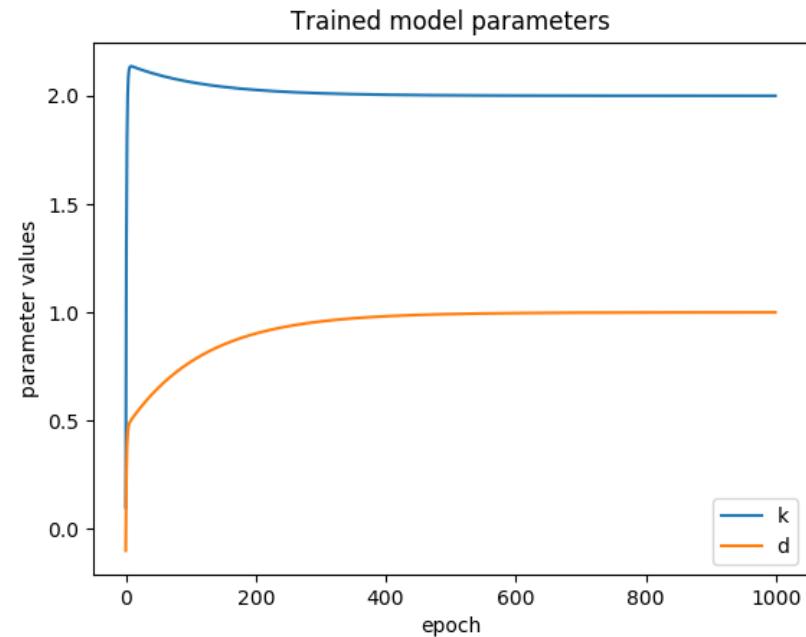
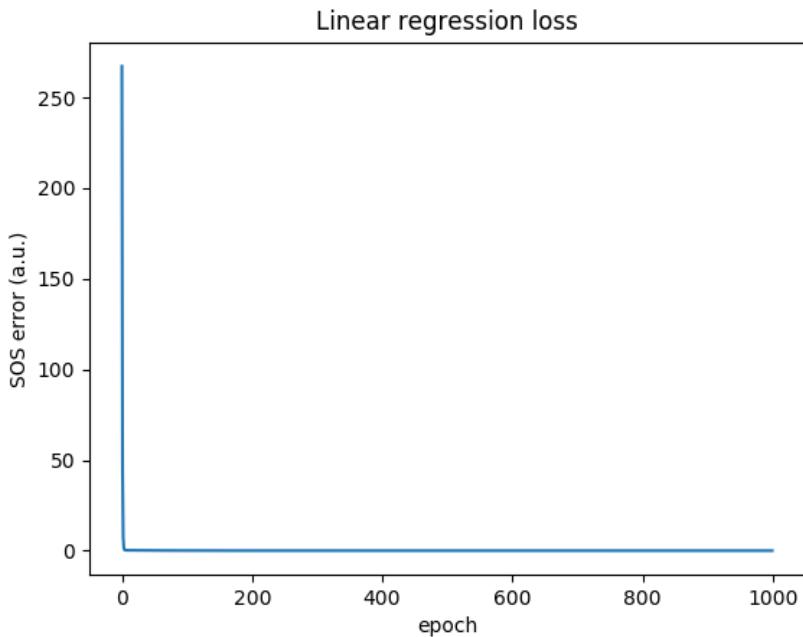
t = time.time()
for ii in range(training_epochs):
    [k_hi_temp, d_hi_temp, loss_hi[ii]] = sess.run([k, d, loss], {x:x_train, y:y_train})
    k_hi[ii] = k_hi_temp.item()
    d_hi[ii] = d_hi_temp.item()
    print('epoch: {}, k={:.3}, b={:.3}'.format(ii+1,k_hi[ii], d_hi[ii], loss_hi[ii]))
    sess.run(optimizer, {x:x_train, y:y_train})

elapsed = time.time() - t
print('Training time: {:.2} s'.format(elapsed))
```



- Computational graph is fixed at this stage
- A session is started
- Data is now fed to the graph

Linear regression in TensorFlow



Example 2: Classification of brain tissue from DTI data

High level neural network modules

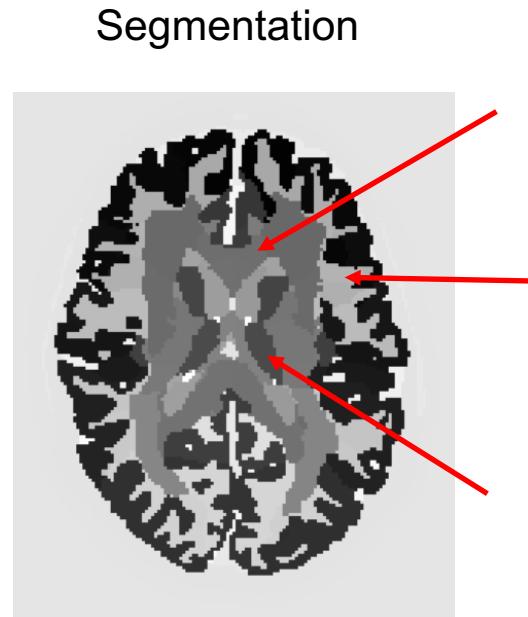
Classification of brain tissue from HCP DTI data



Genu of corpus callusum
(highly aligned WM)

Subcortical WM

Thalamus (GM)



Classification of brain tissue from HCP DTI data

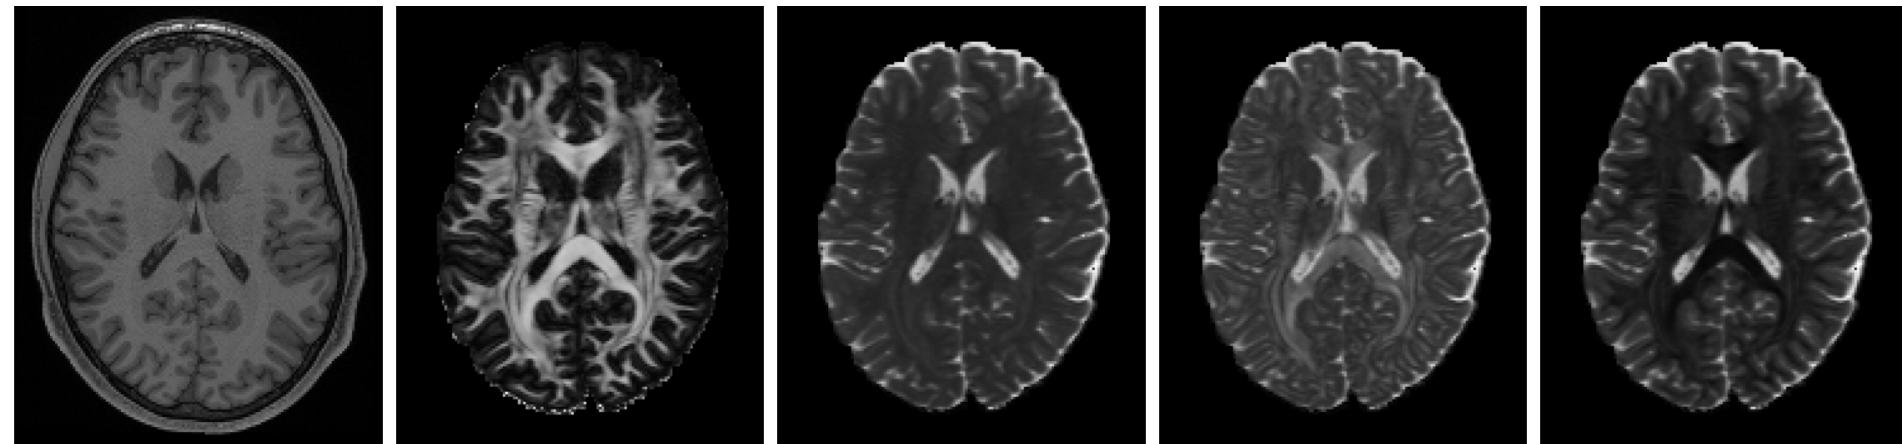
T1w

FA

MD $(\lambda_1 + \lambda_2 + \lambda_3)/3$

AD (λ_1)

RD $(\lambda_2 + \lambda_3)/2$



<https://www.humanconnectome.org/>

ISMRM, June 16th 2018

Structure and inspection of the data

$$\{(x_1, y_1), \dots (x_N, y_N)\}$$

$$x_i = [T1w_i, FA_i, MD_i, AD_i, RD_i]$$

T1w (a.u.)	FA (-)	MD $(\frac{\mu m^2}{ms})$	AD $(\frac{\mu m^2}{ms})$	RD $(\frac{\mu m^2}{ms})$	Class	Class label
898	0.22	1.066592	1.33	0.94	Thalamus	1
1007	0.68	0.39	0.72	0.22	CC	2
867	0.38	0.58	0.82	0.45	Cortical WM	3
...

Normalization!

Training, validation and test set

$$\{(x_1, y_1), \dots (x_N, y_N)\}$$

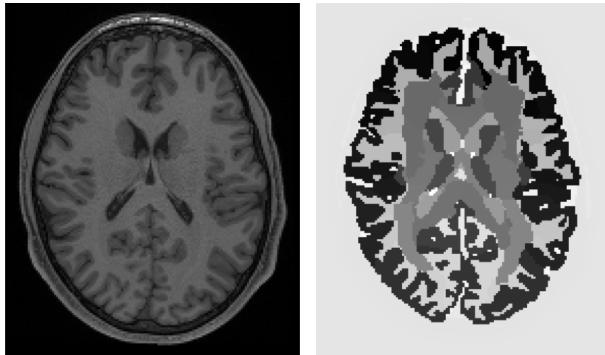
1. Training set: Train the model
2. Validation set: Use as performance criterion to tune training, for hyperparameter selection,...
3. Test set: Don't use during training at all!

Training, validation and test set

1. Training set: 80% randomized samples from 3 subjects
2. Validation set: 20% randomized samples from 3 subjects
3. Test set: Randomized samples from 1 different subject

$$\{(x_1, y_1), \dots (x_N, y_N)\} \quad N = 23700 \quad N_{train} = 14044 \\ N_{val} = 3511 \quad N_{test} = 6145$$

Check balancing of classes



$$N_{thal} = 11662$$

$$N_{cc} = 4582$$

$$N_{subcort-wm} = 7455$$

- Create data duplicates
- Weight samples during training

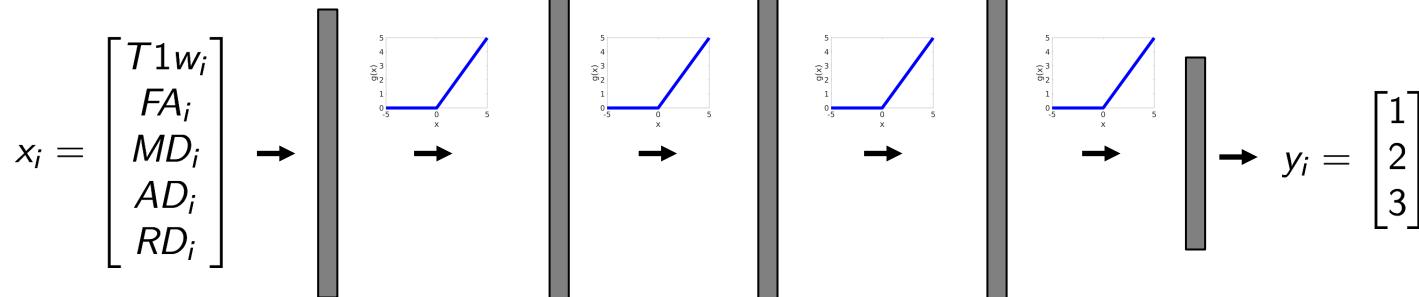
Let's build our first neural network

Multi layer perceptron

- 1 input layer
- 3 hidden layers with 100 elements
- 1 output layer
- Non-linearity: ReLu
- Softmax output

Model Architecture

$$y_i(\mathbf{x}) = g(\mathbf{W}_i \mathbf{x} + \mathbf{b}_i)$$



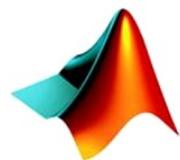
$$(5*100)+100 = 600$$

$$(100*3)+3 = 303$$

Total number of parameters: 31203

$$(100*100)+100 = 10100$$

Let's build our first neural network (for real!)



MATLAB



Neural networks toolbox

PYTORCH



Neural networks package

TensorFlow



Keras

Let's build our first neural network in Matlab

```
%% NN with 3 HIDDEN LAYER NEURONS
nElements = 100;
nLayers = 3;
inputLayer=imageInputLayer([nFeatures,1,1]);
f1=fullyConnectedLayer(nElements);
f2=fullyConnectedLayer(nElements);
f3=fullyConnectedLayer(nElements);
f4=fullyConnectedLayer(nClasses);
s1=softmaxLayer();
outputLayer=classificationLayer();

architecture = [inputLayer; f1; f2; f3; f4; s1; outputLayer];
disp(architecture);
```

```
epochs = 250;
miniBatchSize = 1024;
InitialLearnRate = 0.001;

% Training options: Note that we set the validation patience stopping
% criterion to the number of epochs. This is a stupid thing to do, but we
% want force the training to go to the defined number of epochs so that it
% is consistent with Tensorflow and Pytorch
options = trainingOptions('adam','MaxEpochs',epochs,'InitialLearnRate',InitialLearnRate, ...
    'MiniBatchSize',miniBatchSize,'ExecutionEnvironment','cpu','Plots','training-progress',...
    'ValidationData',{x_val,y_val},'ValidationPatience',epochs);
```

```
%% Train
tic
[net,op] = trainNetwork(x_train,y_train,architecture,options);
toc
```



Kingma and Ba, ICLR 2015

ISMRM, June 16th 2018

Let's build our first neural network in PyTorch



```
%% Define model
nElements = 100
nLayers = 3
model_name = 'dti_FC'
model = torch.nn.Sequential(
    torch.nn.Linear(nFeatures, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nClasses, bias=True),
)
print(model)

%%choose optimizer and loss function
training_epochs = 250
lr = 0.001
batch_size = 1024
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=lr)

%%Create minibatch data loading for training and validation
dataLoader_train = data_utils.TensorDataset(x_train, y_train)
dataLoader_train = data_utils.DataLoader(dataLoader_train, batch_size=batch_size, shuffle=False,num_workers=4)

%% Train model
loss_train = np.zeros(training_epochs)
acc_train = np.zeros(training_epochs)
loss_val = np.zeros(training_epochs)
acc_val = np.zeros(training_epochs)

for epoch in range(training_epochs):
    for local_batch, local_labels in dataLoader_train:
        # feedforward - backpropagation
        optimizer.zero_grad()
        out = model(local_batch)
        loss = criterion(out, local_labels)
        loss.backward()
        optimizer.step()
        loss_train[epoch] = loss.item()
        # Training data accuracy
        [dummy, predicted] = torch.max(out.data, 1)
        acc_train[epoch] = (torch.sum(local_labels==predicted).numpy() / np.size(local_labels.numpy(),0))

        # Validation
        out_val = model(x_val)
        loss = criterion(out_val, y_val)
        loss_val[epoch] = loss.item()
        [dummy, predicted_val] = torch.max(out_val.data, 1)
        acc_val[epoch] = ( torch.sum(y_val==predicted_val).numpy() / setsize_val)

    print ('Epoch {} / {} train loss: {:.3}, train acc: {:.3}, val loss: {:.3}, val acc: {:.3}'.format(epoch+1, training_epochs, loss_train[epoch], acc_train[epoch], loss_val[epoch], acc_val[epoch]))
```

Model architecture

```
## Define model
nElements = 100
nLayers = 3
model_name = 'dti_FC'
model = torch.nn.Sequential(
    torch.nn.Linear(nFeatures, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nClasses, bias=True),
)
print(model)
```



Set up optimizer and data support functions

```
#% Define model
nElements = 100
nLayers = 3
model_name = 'dti_FC'
model = torch.nn.Sequential(
    torch.nn.Linear(nFeatures, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nClasses, bias=True),
)
print(model)

#%choose optimizer and loss function
training_epochs = 250
lr = 0.001
batch_size = 1024
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=lr)

#%Create minibatch data loading for training and validation
dataLoader_train = data_utils.TensorDataset(x_train, y_train)
dataLoader_train = data_utils.DataLoader(dataLoader_train, batch_size=batch_size, shuffle=False,num_workers=4)

#% Train model
loss_train = np.zeros(training_epochs)
acc_train = np.zeros(training_epochs)
loss_val = np.zeros(training_epochs)
acc_val = np.zeros(training_epochs)

for epoch in range(training_epochs):
    for local_batch, local_labels in dataLoader_train:
        # feedforward - backpropagation
        optimizer.zero_grad()
        out = model(local_batch)
        loss = criterion(out, local_labels)
        loss.backward()
        optimizer.step()
        loss_train[epoch] = loss.item()
        # Training data accuracy
        [dummy, predicted] = torch.max(out.data, 1)
        acc_train[epoch] = (torch.sum(local_labels==predicted).numpy() / np.size(local_labels.numpy(),0))

        # Validation
        out_val = model(x_val)
        loss = criterion(out_val, y_val)
        loss_val[epoch] = loss.item()
        [dummy, predicted_val] = torch.max(out_val.data, 1)
        acc_val[epoch] = ( torch.sum(y_val==predicted_val).numpy() / setsize_val)

    print ('Epoch {} / {} train loss: {:.3}, train acc: {:.3}, val loss: {:.3}, val acc: {:.3}'.format(epoch+1, training_epochs, loss_train[epoch], acc_train[epoch], loss_val[epoch], acc_val[epoch]))
```



Kingma and Ba, ICLR 2015

Let's train our network

```
% Define model
nElements = 100
nLayers = 3
model_name = 'dti_FC'
model = torch.nn.Sequential(
    torch.nn.Linear(nFeatures, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nClasses, bias=True),
)
print(model)

#choose optimizer and loss function
training_epochs = 250
lr = 0.001
batch_size = 1024
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=lr)

#Create minibatch data loading for training and validation
dataLoader_train = data_utils.TensorDataset(x_train, y_train)
dataLoader_train = data_utils.DataLoader(dataLoader_train, batch_size=batch_size, shuffle=False,num_workers=4)

# Train model
loss_train = np.zeros(training_epochs)
acc_train = np.zeros(training_epochs)
loss_val = np.zeros(training_epochs)
acc_val = np.zeros(training_epochs)

for epoch in range(training_epochs):
    for local_batch, local_labels in dataLoader_train:
        # feedforward - backpropagation
        optimizer.zero_grad()
        out = model(local_batch)
        loss = criterion(out, local_labels)
        loss.backward()
        optimizer.step()
        loss_train[epoch] = loss.item()
        # Training data accuracy
        [dummy, predicted] = torch.max(out.data, 1)
        acc_train[epoch] = (torch.sum(local_labels==predicted).numpy() / np.size(local_labels.numpy(),0))

        # Validation
        out_val = model(x_val)
        loss = criterion(out_val, y_val)
        loss_val[epoch] = loss.item()
        [dummy, predicted_val] = torch.max(out_val.data, 1)
        acc_val[epoch] = ( torch.sum(y_val==predicted_val).numpy() / setsize_val)

    print ('Epoch {} train loss: {:.3}, train acc: {:.3}, val loss: {:.3}, val acc: {:.3}'.format(epoch+1, training_epochs, loss_train[epoch], acc_train[epoch], loss_val[epoch], acc_val[epoch]))
```



Kingma and Ba, ICLR 2015

Let's train our network



```
## Train model
loss_train = np.zeros(training_epochs)
acc_train = np.zeros(training_epochs)
loss_val = np.zeros(training_epochs)
acc_val = np.zeros(training_epochs)

for epoch in range(training_epochs):
    for local_batch, local_labels in dataloader_train:
        # feedforward - backpropagation
        optimizer.zero_grad()
        out = model(local_batch)
        loss = criterion(out, local_labels)
        loss.backward()
        optimizer.step()
        loss_train[epoch] = loss.item()
        # Training data accuracy
        [dummy, predicted] = torch.max(out.data, 1)
        acc_train[epoch] = (torch.sum(local_labels==predicted).numpy() / np.size(local_labels.numpy(),0))

    # Validation
    out_val = model(x_val)
    loss = criterion(out_val, y_val)
    loss_val[epoch] = loss.item()
    [dummy, predicted_val] = torch.max(out_val.data, 1)
    acc_val[epoch] = ( torch.sum(y_val==predicted_val).numpy() / setsize_val)
```

Let's build our first neural network in TensorFlow

```
## Define model
nElements = 100
nLayers = 3
model = Sequential()
model.add(Dense(nElements, input_dim=nFeatures, activation='relu',name="input_layer"))
model.add(Dense(nElements, activation='relu',name="hidden_layer_01"))
model.add(Dense(nElements, activation='relu',name="hidden_layer_02"))
model.add(Dense(nElements, activation='relu',name="hidden_layer_03"))
model.add(Dense(nClasses, activation='softmax', name="output_softmax"))
model_name = 'dti_FC_{0}layers_{1}elements'.format(nLayers,nElements)

# Compile model
# default adam learning rate: 0.001
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

## Plot model information
model.summary()
with open('./models/{}.txt'.format(model_name), 'w') as fh:
    model.summary(print_fn=lambda x: fh.write(x + '\n'))

plot_model(model, to_file='./models/{}.png'.format(model_name))

## Set up tensorflow
tensorboard_graphdir = './graph_dti/{0}_{1}'.format(time.strftime('%Y-%m-%d_%H-%M-%S'),model_name)
os.makedirs(tensorboard_graphdir)
tbCallBack = keras.callbacks.TensorBoard(log_dir=tensorboard_graphdir, histogram_freq=0, write_graph=True, write_images=True)

## Train model
training_epochs = 250
batch_size=1024

print("Train",model_name,"dti classification")
print("Training data points: {}".format(np.size(y_train,0)))
print("Validation data points: {}".format(np.size(y_val,0)))
print("Test data points: {}".format(np.size(y_test,0)))

training_history=model.fit(x_train, y_train, batch_size=batch_size, validation_data=(x_val,y_val), epochs=training_epochs, verbose=1,callbacks=[tbCallBack])
```



Model architecture

```
## Define model
nElements = 100
nLayers = 3
model = Sequential()
model.add(Dense(nElements, input_dim=nFeatures, activation='relu', name="input_layer"))
model.add(Dense(nElements, activation='relu', name="hidden_layer_01"))
model.add(Dense(nElements, activation='relu', name="hidden_layer_02"))
model.add(Dense(nElements, activation='relu', name="hidden_layer_03"))
model.add(Dense(nClasses, activation='softmax', name="output_softmax"))
model_name = 'dti_FC_{}layers_{}elements'.format(nLayers,nElements)
```

```
# Compile model
# default adam learning rate: 0.001
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```



Let's train our network

```
## Define model
nElements = 100
nLayers = 3
model = Sequential()
model.add(Dense(nElements, input_dim=nFeatures, activation='relu',name="input_layer"))
model.add(Dense(nElements, activation='relu',name="hidden_layer_01"))
model.add(Dense(nElements, activation='relu',name="hidden_layer_02"))
model.add(Dense(nElements, activation='relu',name="hidden_layer_03"))
model.add(Dense(nClasses, activation='softmax', name="output_softmax"))
model_name = 'dti_FC_{0}layers_{1}elements'.format(nLayers,nElements)

# Compile model
# default adam learning rate: 0.001
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

## Plot model information
model.summary()
with open('./models/{}.txt'.format(model_name), 'w') as fh:
    model.summary(print_fn=lambda x: fh.write(x + '\n'))

plot_model(model, to_file='./models/{}.png'.format(model_name))

## Set up tensorflow
tensorboard_graphdir = './graph_dti/{0}_{1}'.format(time.strftime('%Y-%m-%d_%H-%M-%S'),model_name)
os.makedirs(tensorboard_graphdir)
tbCallBack = keras.callbacks.TensorBoard(log_dir=tensorboard_graphdir, histogram_freq=0, write_graph=True, write_images=True)

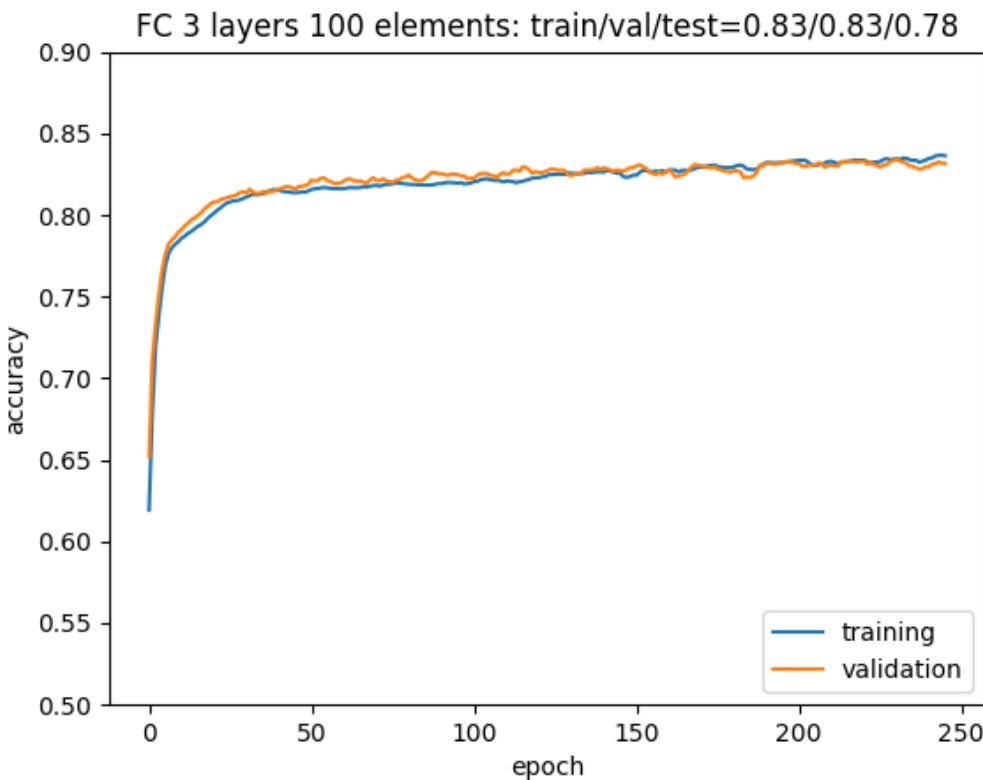
## Train model
training_epochs = 250
batch_size=1024

print("Train",model_name,"dti classification")
print("Training data points: {}".format(np.size(y_train,0)))
print("Validation data points: {}".format(np.size(y_val,0)))
print("Test data points: {}".format(np.size(y_test,0)))

training_history=model.fit(x_train, y_train, batch_size=batch_size, validation_data=(x_val,y_val), epochs=training_epochs, verbose=1,callbacks=[tbCallBack])
```



Results



- Higher model complexity?
- Inspect misclassified samples
- Trends for different label categories?
- Experiment with features
- Don't use test performance to adjust model architecture!

Example 3: Classification of image quality of
accelerated reconstructions

Convolutional Neural Networks (CNNs),
model complexity

Fully sampled vs 4 times PI-CS accelerated

Fully sampled reference



PI-CS (TGV) R=4



Fully sampled vs 4 times PI-CS accelerated

Fully sampled reference



PI-CS (TGV) R=4



Fully sampled vs 4 times PI-CS accelerated

Fully sampled reference



PI-CS (TGV) R=4



Fully sampled vs 4 times PI-CS accelerated

Fully sampled reference



PI-CS (TGV) R=4



Training, validation and test set

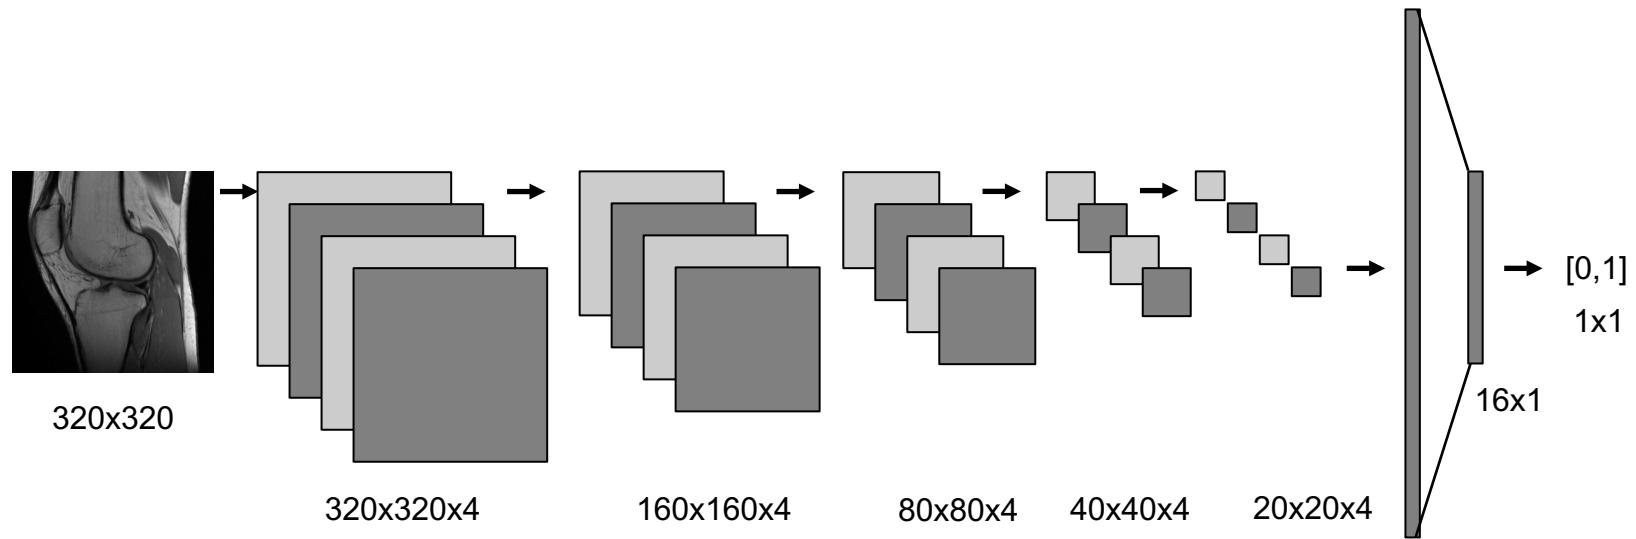
1. Training set: 2/3 randomized samples
2. Validation set: 1/3 randomized samples
3. No separate test set: Experiment is left to the audience :-)

$$N = 520$$

$$\{(x_1, y_1), \dots (x_N, y_N)\} \quad N_{train} = 390 \quad N_{ref} = 260$$

$$N_{val} = 130 \quad N_{tgv} = 260$$

Our first convolutional neural network (CNN)



4 3x3x4 convolutional layers, Relu activation

2D MaxPooling

1 fully connected layer with 16 elements

Sigmoid Output

Total number of parameters: 26117

Our first convolutional neural network (CNN)

```
## Define model architecture
model = Sequential()
model.add(Convolution2D(4, (3, 3), activation='relu', padding='same', input_shape=(nR,nC,1)))
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Convolution2D(4, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Convolution2D(4, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Convolution2D(4, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2,2)))

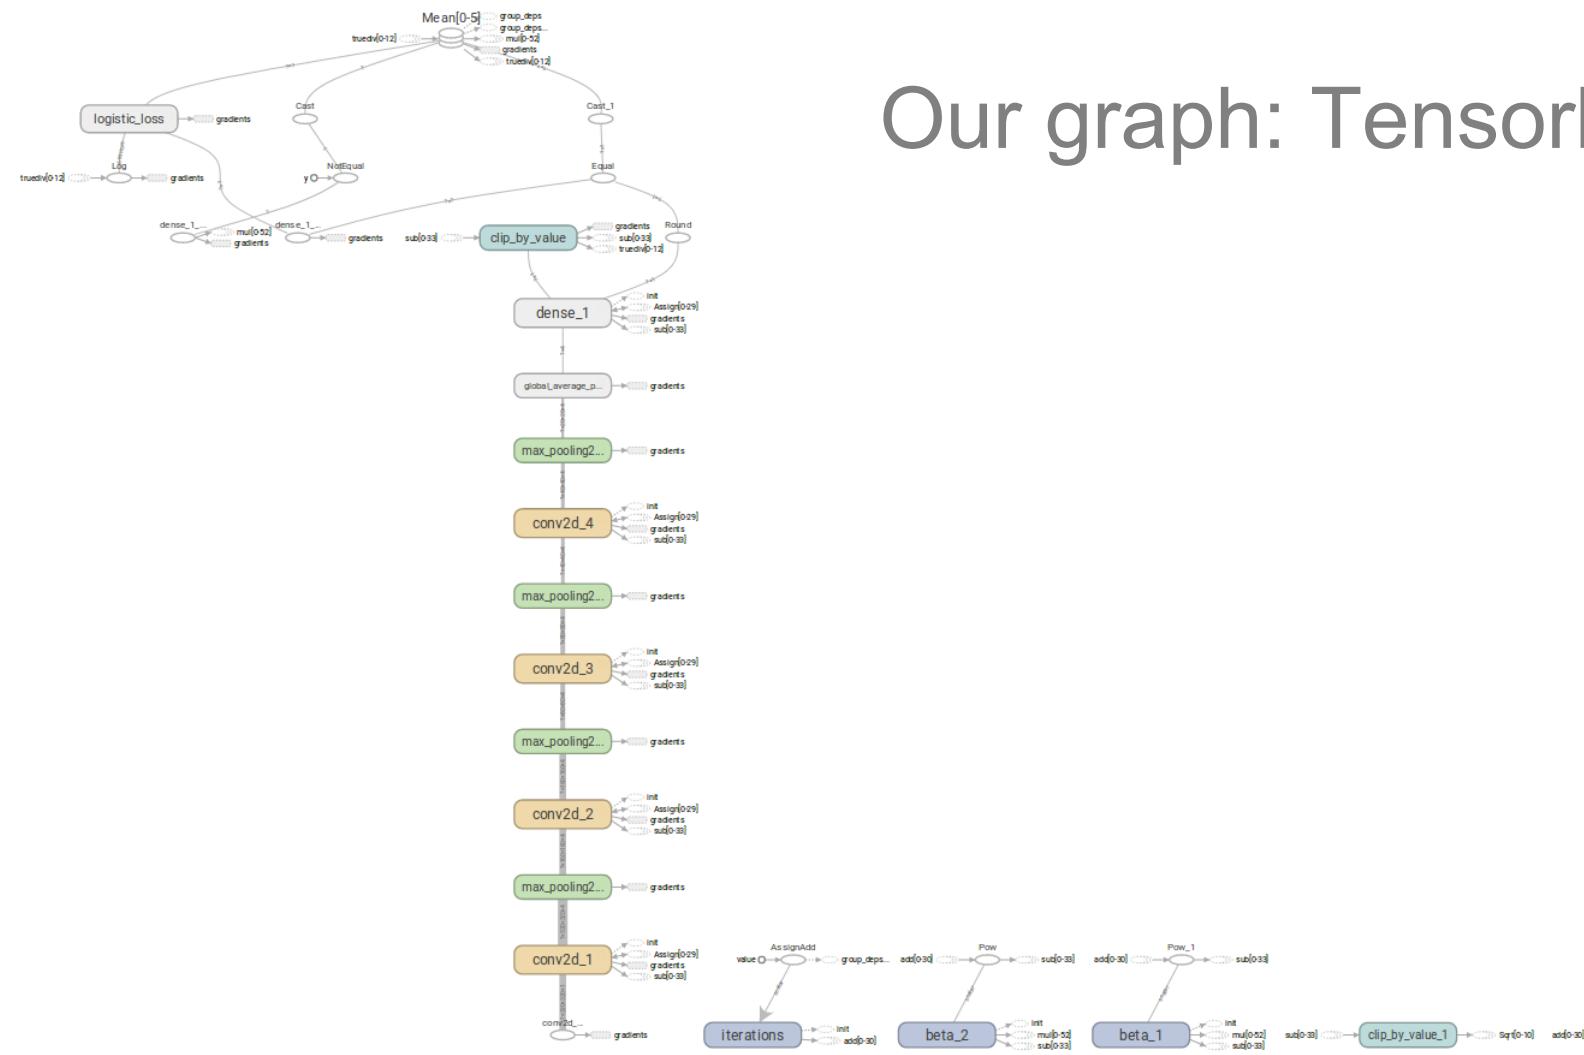
# Global Average Pooling
#model_name = 'CNN1layers_global_avg';
#model_name = 'CNN4layers_global_avg';
#model.add(GlobalAveragePooling2D())

# Fully connected and dense layer
model_name = 'CNN4layers_FC';
model.add(Flatten())
model.add(Dense(16, activation='relu'))

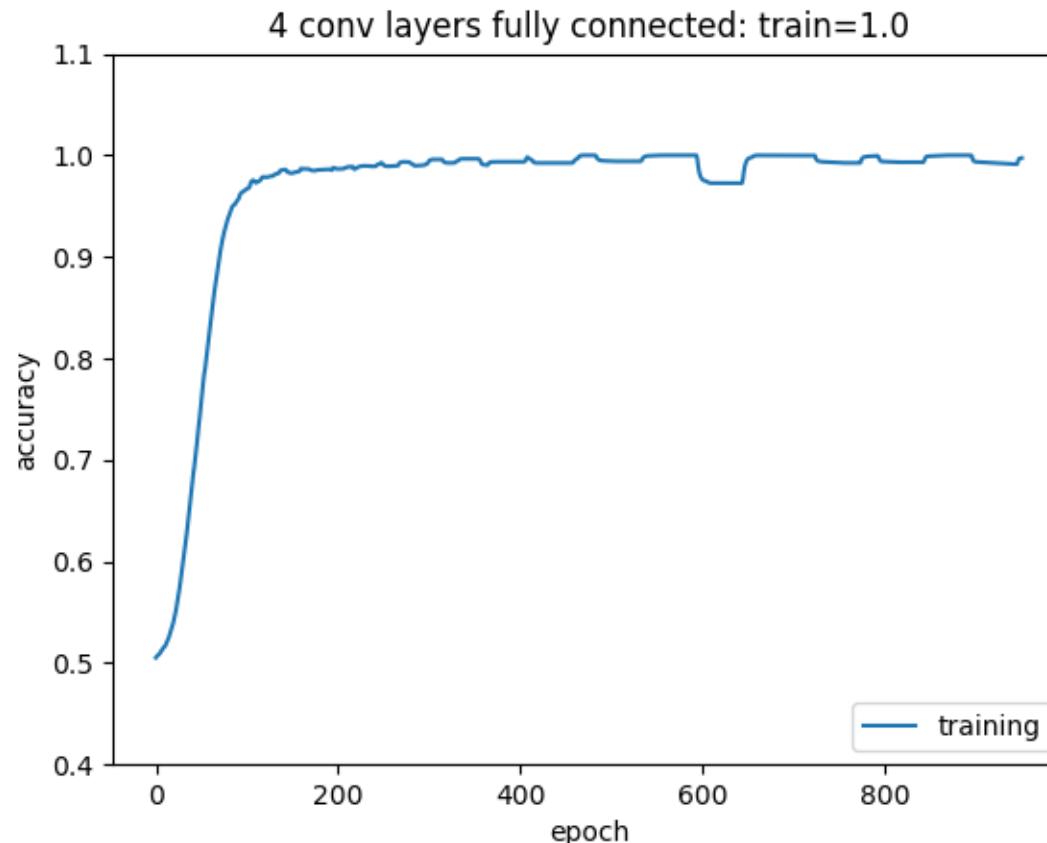
# Output layer
model.add(Dense(1, activation='sigmoid'))
```



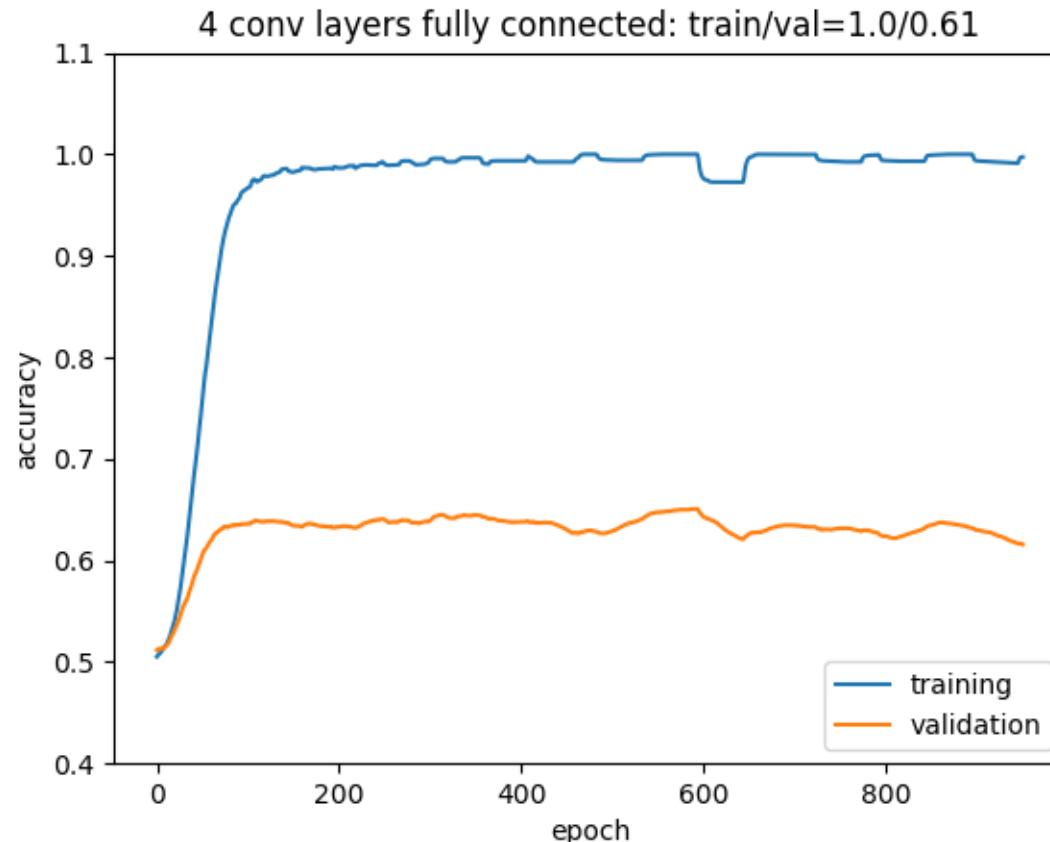
Our graph: Tensorboard



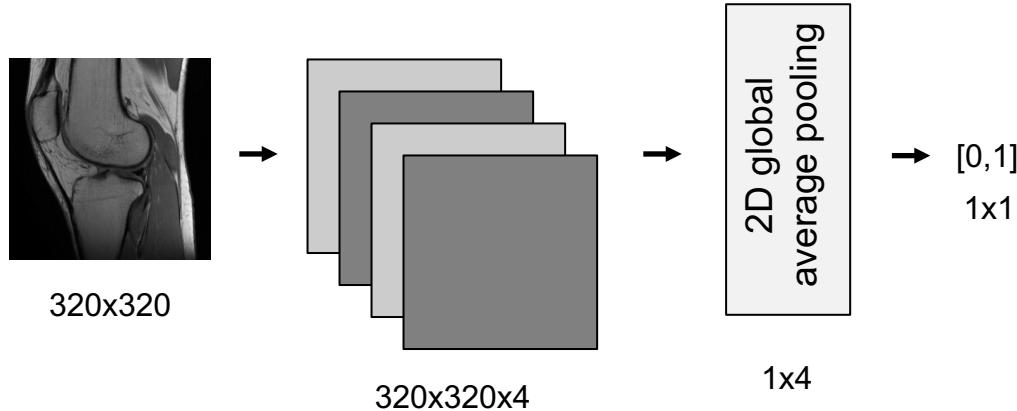
Results



Results



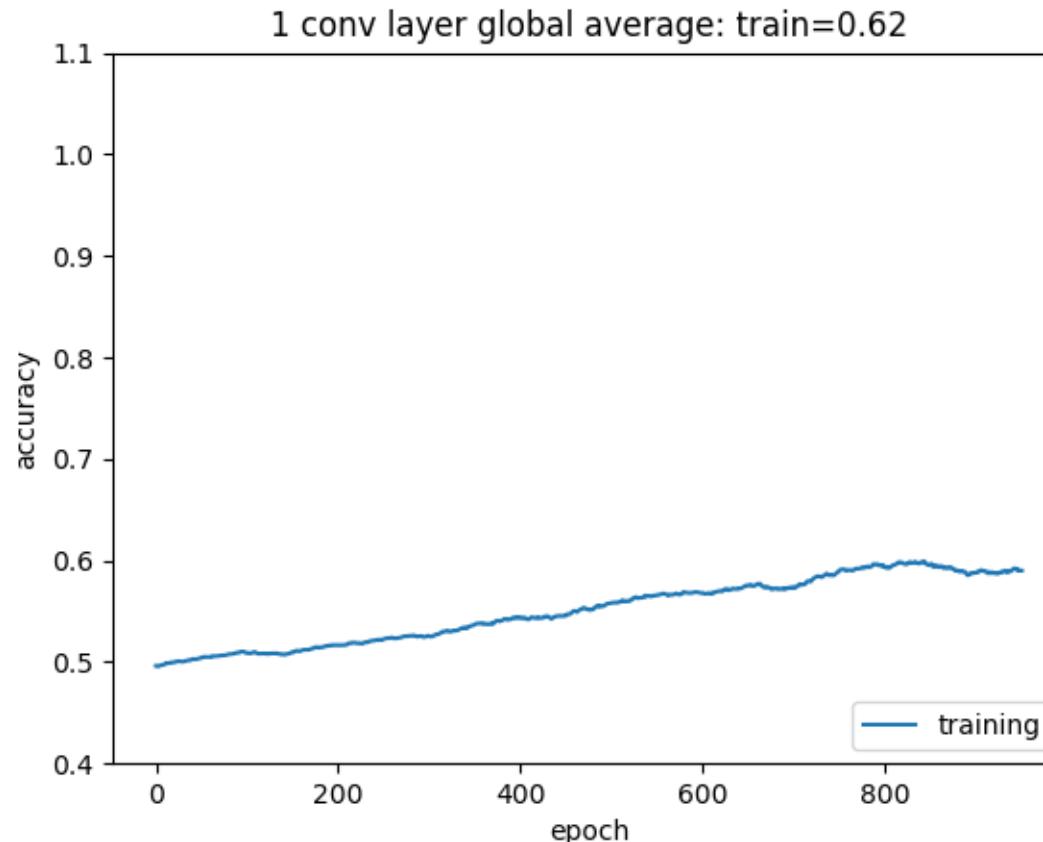
Reduce complexity of our model



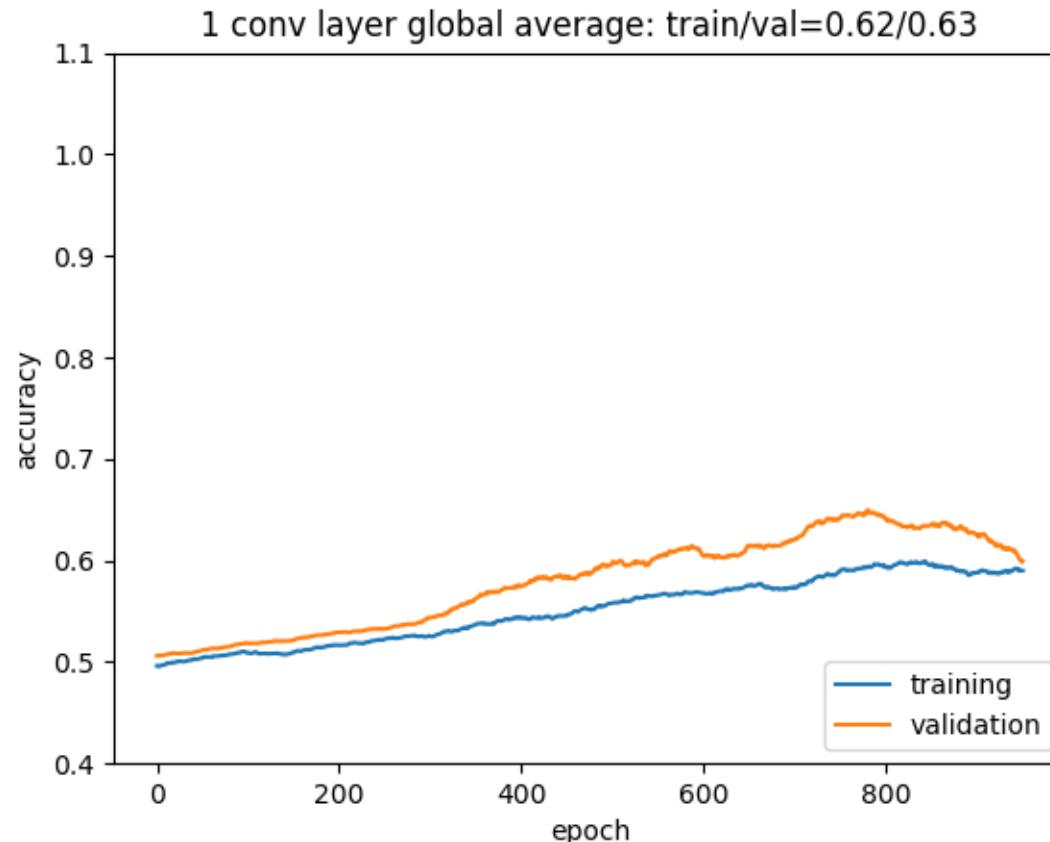
1 $3 \times 3 \times 4$ convolutional layer, Relu activation
2D Global Average Pooling
Sigmoid Output

Total number of parameters: 45

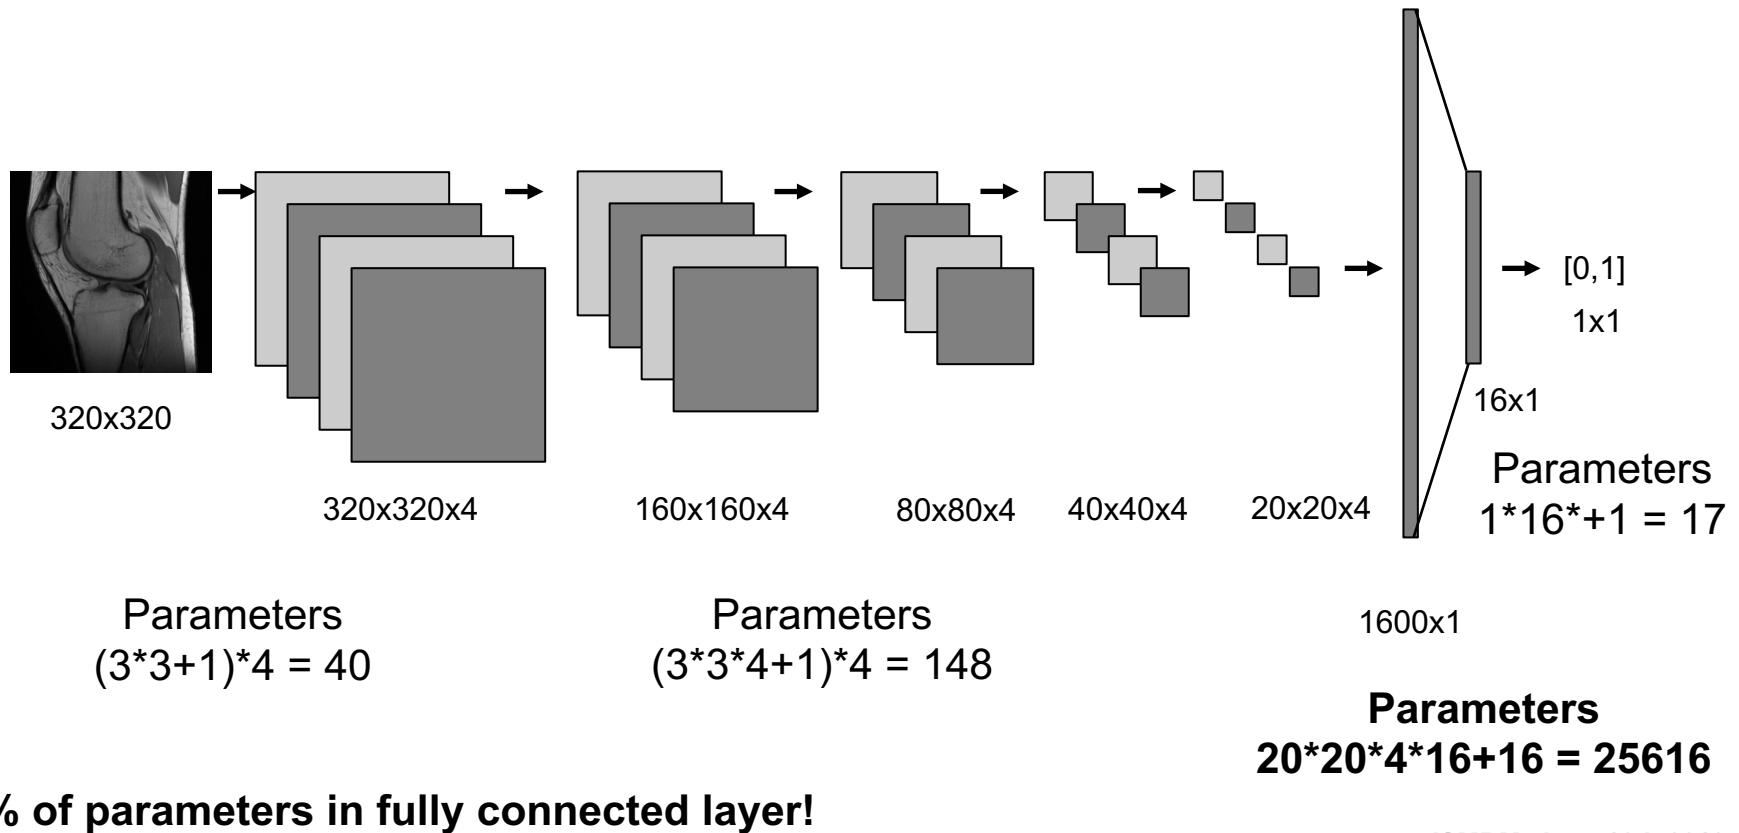
Results



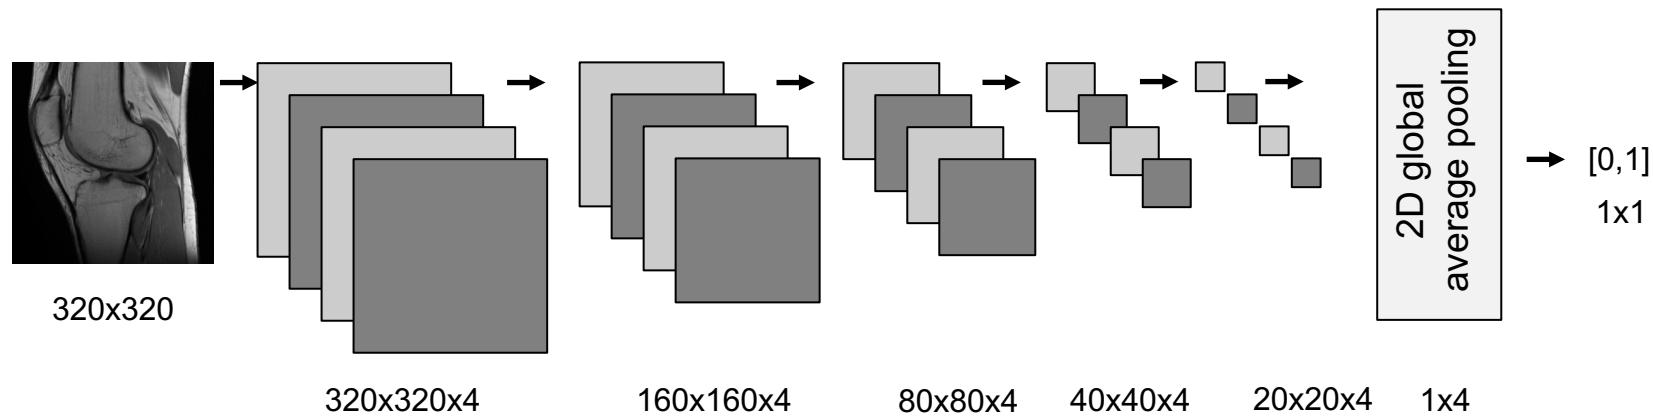
Results



A closer look at the parameter distribution of the first model



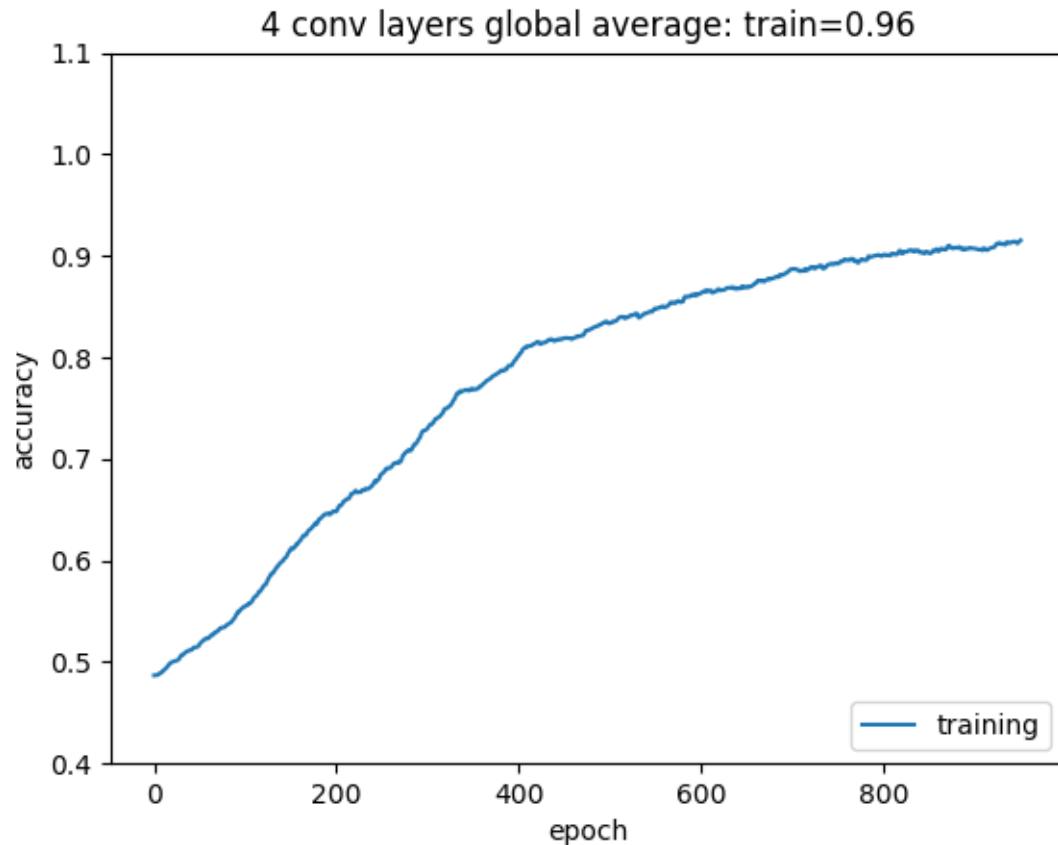
Replace FC layer with global average pooling



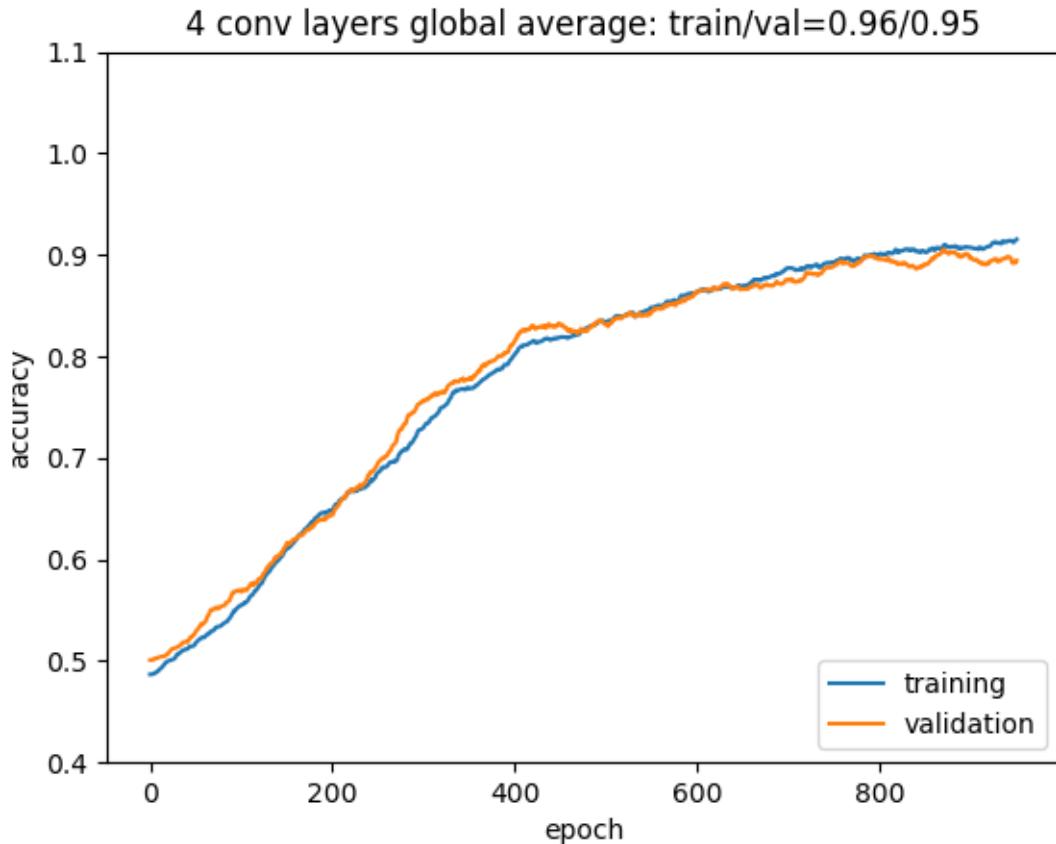
4 3x3x4 convolutional layers, Relu activation
2D MaxPooling
2D Global Average Pooling

Total number of parameters: 489

Results



Results



- Use separate test set
- Inspect misclassified samples
- Number of conv layers, filter size,...
- Use higher model complexity combined with regularization (e.g. dropout)

Srivastava et al., JMLR (2014)

ISMRM, June 16th 2018

Summary

- Examples for linear regression, MLPs, CNNs
- Software frameworks
- High level view on generalization
- Model complexity, overfitting vs. underfitting
- Model parameters: Convolutional vs fully connected layers

Upcoming workshops



ISMRM Workshop on Machine Learning, Part II

25-28 October 2018

Capital Hilton, Washington, D.C., USA

Machine Learning for Medical Image
Reconstruction (MLMIR)

MICCAI2018
Granada
SPAIN

- MICCAI paper results released: 22nd May 2018
- Paper submission deadline: 11th June 2018
- Decisions to authors: 2nd July 2018
- Camera-ready submission: 17th July 2018
- Workshop: 16th September 2018 (am), Granada, Spain



ML solves the problems the world really cares about!

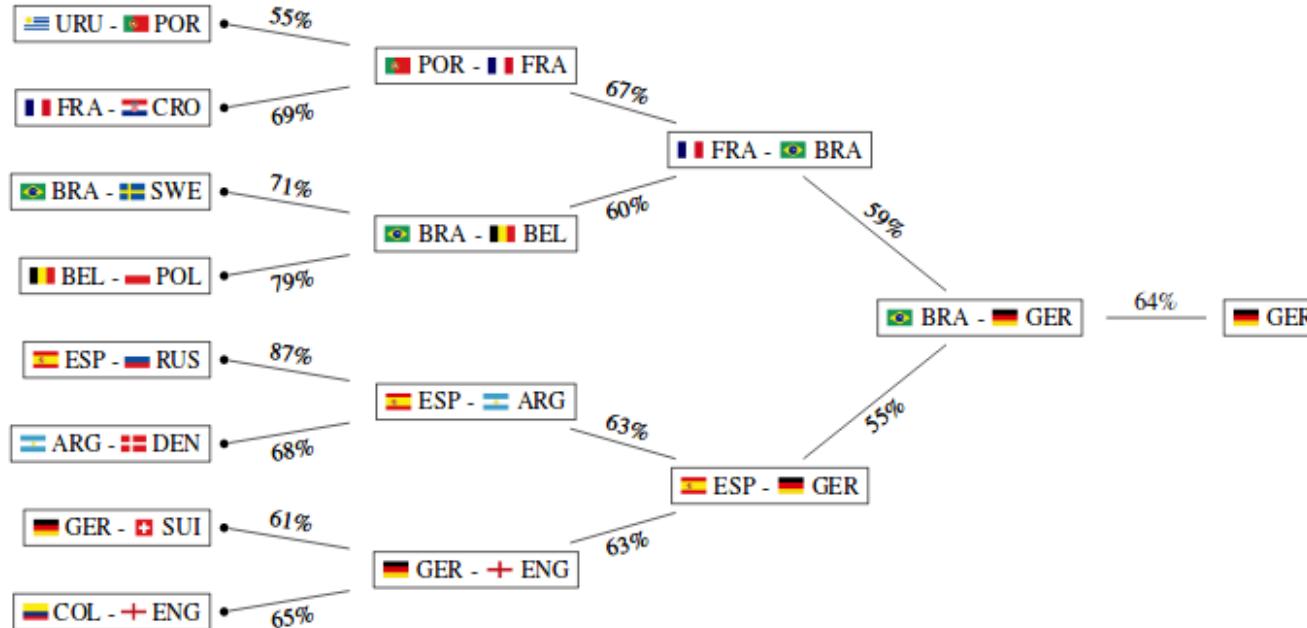


Figure 5: Most probable course of the knockout stage together with corresponding probabilities for the FIFA World Cup 2018 based on 100,000 simulation runs.

Acknowledgements

- NYU
 - Mary Bruno
 - Joe Katsnelson
 - Matt Muckley
 - Michael Recht
 - Dan Sodickson
- Siemens healthineers
 - Matthias Fenchel
- Grant Support
 - NIH P41 EB017183
- TU Graz
 - Kerstin Hammernik
 - Erich Kobler
 - Thomas Pock

Thank you!

https://github.com/FlorianKnoll/ISMRM2018_Educational_DeepLearning

